

Audio Processing Pipeline Diagrams

PURPOSE: Visual reference for the complete audio processing chain from API enqueue through queue management, decoder chains, buffer management, mixer, and audio output.

AUDIENCE: Developers, architects, and technical documentation readers

RELATED DOCUMENTATION:

- [SPEC016 Decoder Buffer Design](#)
- [SPEC013 Single Stream Playback](#)
- [SPEC014 Single Stream Design](#)
- [SPEC002 Crossfade Design](#)

Table of Contents

1. [High-Level Overview \(Mermaid Flowchart\)](#)
 2. [Interaction Flow \(Mermaid Sequence Diagram\)](#)
 3. [Buffer Lifecycle \(Mermaid State Diagram\)](#)
 4. [Comprehensive Reference \(ASCII Diagram\)](#)
 5. [DBD-PARAM Parameter Mapping](#)
-

1. High-Level Overview

Linear Pipeline Flow

```
flowchart TD
    Start[HTTP POST /playback/enqueue] --> API[API Handler<br/>handlers.rs:311]
    API --> Queue[QueueManager.enqueue<br/>queue_manager.rs:345]
    Queue --> Chain[Assign Chain Index 0-11<br/>DBD-LIFECYCLE-040]
    Chain --> Submit[DecoderWorker.submit<br/>decoder_worker.rs:145]
    Submit --> Create[Create DecoderChain<br/>decoder_chain.rs:104]

    Create --> Decode[StreamingDecoder<br/>Read PCM from file]
    Decode --> Resample[StatefulResampler<br/>44.1kHz]
    Resample --> Fade[Fader<br/>Apply curves]
    Fade --> Buffer[PlayoutRingBuffer<br/>via BufferManager]

    Buffer --> Ready{Buffer Ready?<br/>threshold reached}
    Ready -->|Yes| Event[BufferEvent::ReadyForStart<br/>PERF-POLL-010]
    Ready -->|No| Decode

    Event --> Mixer[CrossfadeMixer<br/>pipeline/mixer.rs]
    Mixer --> Ring[AudioRingBuffer<br/>lock-free SPSC]
    Ring --> Callback[Audio Callback<br/>Real-time thread]
    Callback --> Output[CPAL Stream<br/>Speakers]

    style Start fill:#e1f5ff
    style Output fill:#e8f5e9
```

```
style Event fill:#fff9c4
style Mixer fill:#f3e5f5
```

Component Architecture

```
graph TB
    subgraph "API Layer"
        API[API Handler<br/>POST /playback/enqueue]
    end

    subgraph "Queue Management"
        QM[QueueManager<br/>current/next/queued]
        DB[(SQLite DB<br/>queue table)]
        QM <--> DB
    end

    subgraph "PlaybackEngine Coordination"
        Engine[PlaybackEngine<br/>engine.rs]
        PBLoop[Playback Loop]
        BEH[Buffer Event Handler]
        PEH[Position Event Handler]
        Engine --> PBLoop
        Engine --> BEH
        Engine --> PEH
    end

    subgraph "Decoder Worker (Single-threaded)"
        DW[DecoderWorker<br/>decoder_worker.rs]
        PQ[Priority Queue<br/>Immediate > Next > Prefetch]
        Active[Active Chains<br/>HashMap]
        DW --> PQ
        DW --> Active
    end

    subgraph "DecoderChain [0-11] - Per Passage Pipeline"
        DC[DecoderChain]
        SD[StreamingDecoder<br/>symphonia]
        SR[StatefulResampler<br/>rubato]
        F[Fader<br/>fade curves]
        DC --> SD
        SD --> SR
        SR --> F
    end

    subgraph "Buffer Management"
        BM[BufferManager<br/>buffer_manager.rs]
        PRB1[PlayoutRingBuffer 0]
        PRB2[PlayoutRingBuffer 1]
        PRBn[PlayoutRingBuffer N]
        BM --> PRB1
        BM --> PRB2
    end
```

```

    BM --> PRBn
end

subgraph "Mixing & Output"
    MX[CrossfadeMixer<br/>pipeline/mixer.rs]
    MT[Mixer Thread<br/>Graduated filling]
    ARB[AudioRingBuffer<br/>lock-free output]
    AC[Audio Callback<br/>real-time thread]
    MX --> MT
    MT --> ARB
    ARB --> AC
end

subgraph "Audio Device"
    CPAL[CPAL Stream<br/>44.1kHz stereo]
    DEV[Audio Device<br/>Speakers]
    CPAL --> DEV
end

API --> QM
QM --> Engine
Engine --> DW
DW --> DC
F --> BM
BM -.ReadyForStart.-> BEH
BEH --> MX
PRB1 --> MX
PRB2 --> MX
PRBn --> MX
AC --> CPAL

style API fill:#e1f5ff
style DW fill:#fff3e0
style DC fill:#f3e5f5
style BM fill:#e8f5e9
style MX fill:#fce4ec
style CPAL fill:#e0f2f1

```

2. Interaction Flow

Sequence Diagram: Enqueue to Playback

```

sequenceDiagram
    participant User as HTTP Client
    participant API as API Handler
    participant Queue as QueueManager
    participant Engine as PlaybackEngine
    participant Worker as DecoderWorker
    participant Chain as DecoderChain
    participant BufMgr as BufferManager

```

```

participant Mixer as CrossfadeMixer
participant Output as AudioOutput

User->>API: POST /playback/enqueue<br/>{file_path}
API->>Queue: enqueue(file_path)
Queue->>Queue: Add to current/next/queued
Queue-->>API: queue_entry_id (UUID)

API->>Engine: assign_chain(queue_entry_id)
Engine->>Engine: Allocate chain index (0-11)
Note over Engine: DBD-LIFECYCLE-030<br/>Lowest-numbered allocation

API->>Worker: submit(DecodeRequest)
Note over Worker: Priority: Immediate/Next/Prefetch
Worker->>Chain: create_chain(queue_entry_id, passage)

Chain->>BufMgr: allocate_buffer(queue_entry_id)
BufMgr-->>Chain: PlayoutRingBuffer
Note over BufMgr: State: Empty<br/>DBD-BUF-020

loop Every ~1 second chunk
    Chain->>Chain: decode_chunk()
    Chain->>Chain: resample(→ 44.1kHz)
    Chain->>Chain: apply_fades()
    Chain->>BufMgr: push_samples(samples)

    alt Buffer reaches threshold
        BufMgr->>BufMgr: Empty → Filling → Ready
        BufMgr-->>Engine: BufferEvent::ReadyForStart
        Note over BufMgr: DBD-PARAM-088<br/>mixer_min_start_level<br/>(default:
22050 samples)
    end
end

Engine->>Mixer: start_passage(queue_entry_id)

loop Continuous playback
    Mixer->>BufMgr: pop_frames(queue_entry_id)
    BufMgr-->>Mixer: AudioFrame
    Mixer->>Mixer: Apply volume
    Mixer->>Mixer: Apply crossfade (if transitioning)
    Mixer->>Output: push to AudioRingBuffer
end

loop Audio callback (~50ms intervals)
    Output->>Output: Pop from AudioRingBuffer
    Output->>Output: Write to device buffer
    Output->>User: Audio to speakers
end

API-->>User: 200 OK<br/>{queue_entry_id}

```

Event-Driven Architecture

```
sequenceDiagram
    participant BufMgr as BufferManager
    participant BufEvt as Buffer Event Channel
    participant PosEvt as Position Event Channel
    participant Engine as Engine Tasks
    participant Mixer as CrossfadeMixer
    participant State as SharedState
    participant SSE as SSE Clients

    Note over BufMgr,BufEvt: Buffer Event Flow
    BufMgr->>BufEvt: BufferEvent::ReadyForStart
    BufEvt->>Engine: buffer_event_handler receives
    Engine->>Mixer: Trigger passage start

    Note over Mixer,PosEvt: Position Event Flow
    loop Every ~1 second
        Mixer->>PosEvt: PlaybackEvent<br/>{queue_entry_id, frame_pos}
        PosEvt->>Engine: position_event_handler receives
        Engine->>State: Update current passage
    end

    Note over State,SSE: SSE Broadcast
    State->>SSE: PlaybackPosition event
    State->>SSE: QueueChanged event
    State->>SSE: BufferChainStatus event
```

3. Buffer Lifecycle

Buffer State Machine

```
stateDiagram-v2
    [*] --> Empty: allocate_buffer()<br/>DBD-BUF-020

    Empty --> Filling: First sample written<br/>decoder writes to buffer

    Filling --> Ready: Threshold reached<br/>DBD-PARAM-088<br/>mixer_min_start_level<br/>(default: 22,050 samples)

    Ready --> Playing: Mixer starts consuming<br/>mixer.start_passage()

    Playing --> Playing: Continuous consumption<br/>mixer.pop_frames()

    Playing --> Finished: Buffer exhausted<br/>All samples consumed

    Finished --> [*]: Buffer released<br/>Chain deallocated

    note right of Ready
```

```

    Emits BufferEvent::ReadyForStart
    (PERF-POLL-010)
    Instant startup notification
end note

note right of Playing
    Buffer may continue filling
    while mixer is consuming
    (simultaneous read/write)
end note

note right of Finished
    DBD-LIFECYCLE-020
    Chain returns to pool
    Lowest-numbered allocation
end note

```

Decoder Pause/Resume State Machine

```

stateDiagram-v2
    [*] --> Decoding: Start chunk processing

    Decoding --> CheckBuffer: Chunk complete

    CheckBuffer --> Decoding: free_space > threshold

    CheckBuffer --> Paused: free_space ≤ headroom<br/>DBD-PARAM-080<br/>(4,410 samples)

    Paused --> CheckResume: Mixer consuming

    CheckResume --> Paused: free_space < resume_threshold

    CheckResume --> Decoding: free_space ≥ resume_threshold<br/>DBD-PARAM-085 + DBD-PARAM-080<br/>(44,100 + 4,410 = 48,510 samples)

    note right of Paused
        Hysteresis prevents oscillation
        Gap = 44,100 samples (1.0 second)
        DBD-PARAM-085
    end note

    note right of Decoding
        Processes ~1 second chunks
        DBD-PARAM-065
        (25,000 samples output)
    end note

```

Mixer Mode State Machine

```

stateDiagram-v2
    [*] --> None: Idle (no passages)

    None --> Playing: Buffer ready + play state

    Playing --> Crossfading: At fade_out_point_ticks<br/>Next passage ready

    Crossfading --> Playing: Crossfade complete<br/>Transition to new passage

    Playing --> PausedDecay: engine.pause() called

    PausedDecay --> PausedDecay: Apply decay_factor<br/>DBD-PARAM-090<br/>(0.96875
per sample)

    PausedDecay --> PausedSilence: level < decay_floor<br/>DBD-PARAM-100<br/>
(0.0001778)

    PausedSilence --> PausedSilence: Output zeros

    PausedDecay --> Playing: engine.play() called<br/>Linear fade-in (500ms)

    PausedSilence --> Playing: engine.play() called<br/>Linear fade-in (500ms)

    Playing --> None: Queue empty

    Crossfading --> None: Error condition

    note right of Crossfading
        Reads from TWO buffers simultaneously
        old_frame * fade_out_weight +
        new_frame * fade_in_weight
    end note

    note right of PausedDecay
        Exponential decay reduces "pop"
        decay^n creates smooth fade to zero
    end note

```

4. Comprehensive Reference

Complete ASCII Pipeline Diagram

```

WKMP-AP AUDIO PROCESSING PIPELINE
(API → Queue → Decode → Mix → Output)

```

1. API REQUEST (handlers.rs:311-366)

```

HTTP POST /playback/enqueue
|
| EnqueueRequest { file_path }
▼
[API Handler: enqueue_passage]
|
| PathBuf::from(file_path)
▼
engine.enqueue_file(file_path)
|
| Emits: QueueChanged, QueueStateUpdate (SSE)
↳ Returns queue_entry_id (UUID)

```

2. QUEUE MANAGEMENT (engine.rs + queue_manager.rs:90-451)

```

engine.enqueue_file()
|
| ↳ [Load/Create Passage from DB]
|   • passages::get_passage_with_timing()
|   • OR passages::create_ephemeral_passage()
|
| ↳ [QueueManager::enqueue]
|   • State: current → next → queued[]
|   • Updates play_order in database
|
| ↳ [Chain Assignment - DBD-LIFECYCLE-040]
|   • Allocate chain_index (0-11 from pool)
|   • chain_assignments[queue_entry_id] = chain_index
|   • DBD-PARAM-050: maximum_decode_streams = 12
|
| ↳ [Submit Decode Request]
|   • decoder_worker.submit(queue_entry_id, passage, priority)
|   • priority = Immediate (current), Next, or Prefetch

```

3. DECODER WORKER (decoder_worker.rs:1-200)

```

DecoderWorker receives DecodeRequest
|
| Priority queue: Immediate > Next > Prefetch
| DBD-PARAM-060: Check priority every 5000ms
▼
[Worker Loop - Single-threaded serial decoding]
|
| ↳ Pending Request?
|   ↳ Create DecoderChain

```


- |
 - └→ Active Chains?
 - └→ Process one chunk (~1 second audio)
 - DBD-PARAM-065: 25,000 samples/chunk (at 44.1kHz)
- |
 - └→ Yielded Chains?
 - └→ Retry if buffer has space

4. DECODER CHAIN (pipeline/decoder_chain.rs:1-200)

DecoderChain::new(queue_entry_id, chain_index, passage)

- |
 - └→ [StreamingDecoder]
 - symphonia-based audio decoder
 - Reads from file (start_ms → end_ms)
 - Outputs PCM samples at source rate
 - |
 - └→ [StatefulResampler]
 - rubato-based resampler
 - Source rate → 44.1kHz (TARGET_SAMPLE_RATE)
 - DBD-PARAM-020: working_sample_rate = 44,100 Hz
 - Maintains state for streaming
 - |
 - └→ [Fader]
 - Applies fade-in curve (at fade_in_point)
 - Applies fade-out curve (at fade_out_point)
 - 5 curve types: linear, exponential, etc.
 - |
 - └→ [Buffer Allocation]
 - buffer_manager.allocate_buffer(queue_entry_id)
 - Creates PlayoutRingBuffer
 - DBD-PARAM-070: playout_ringbuffer_size = 661,941 samples (15.01s)
 - DBD-PARAM-080: playout_ringbuffer_headroom = 4,410 samples (0.1s)
 - DBD-PARAM-085: decoder_resume_hysteresis = 44,100 samples (1.0s)

DecoderChain::process_chunk()

- |
 - [DBD-DEC-110] Process ~1 second chunks
 - |
 - └→ STEP 1: Decode
 - decoder.decode_chunk() → Vec<f32> (PCM samples)
 - |
 - └→ STEP 2: Resample
 - resampler.process(samples) → Vec<f32> (44.1kHz)
 - DBD-PARAM-065: 25,000 samples output per chunk
 - |
 - └→ STEP 3: Fade
 - fader.apply_fades(samples) → Vec<f32> (with curves)

```

└─┐ STEP 4: Push to Buffer
    |
    | buffer_manager.push_samples(queue_entry_id, samples)
    |
    └─┐ Result:
        • ChunkProcessResult::Processed → Continue
        • ChunkProcessResult::BufferFull → Yield
          (free_space ≤ playout_ringbuffer_headroom)
        • ChunkProcessResult::Finished → Done

```

5. BUFFER MANAGER (buffer_manager.rs:1-300)

BufferManager maintains HashMap<Uuid, ManagedBuffer>

```

└─┐ [ManagedBuffer]
    |
    | • PlayoutRingBuffer (lock-free ring buffer)
    | • BufferMetadata (state machine)
    | • States: Empty → Filling → Ready → Playing → Finished
    |
    └─┐ [State Transitions - DBD-BUF-020 through DBD-BUF-060]
        |
        | └─┐ Empty → Filling (first sample written)
            |
            | └─┐ Filling → Ready (threshold reached)
                |
                | • DBD-PARAM-088: mixer_min_start_level = 22,050 samples (0.5s)
                | • For first passage: may use lower threshold (500ms)
                |   └─┐ Emits BufferEvent::ReadyForStart (PERF-POLL-010)
                    |
                    | └─┐ Ready → Playing (mixer starts consuming)
                        |
                        | └─┐ Playing → Finished (all samples consumed)
                            |
                            |
                            └─┐ [Decoder Pause Logic]
                                |
                                | • Pause when: free_space ≤ playout_ringbuffer_headroom (4,410)
                                | • Resume when: free_space ≥ resume_threshold (48,510)
                                | • Resume threshold = DBD-PARAM-085 + DBD-PARAM-080
                                | •               = 44,100 + 4,410 = 48,510 samples
                                | • Hysteresis gap = 44,100 samples (1.0 second)
                                |
                                └─┐ [PlayoutRingBuffer Operations]
                                    |
                                    | • push_samples() - Decoder writes here
                                    | • pop_frames() - Mixer reads here
                                    | • Lock-free ring buffer with capacity tracking
                                    | • DBD-PARAM-070: Capacity = 661,941 samples (15.01s @ 44.1kHz)

```

6. PLAYBACK ENGINE COORDINATION (engine.rs:346-530)

`engine.start()` spawns multiple tasks:

A. Playback Loop (`playback_loop`)

- Monitors queue state
- Triggers passage transitions
- Calculates decode priorities
- Submits decode requests

|

B. Buffer Event Handler
(`buffer_event_handler`)

- Receives `BufferEvent::ReadyForStart`
- Triggers mixer to start passage
- PERF-POLL-010: Instant startup

|

C. Position Event Handler
(`position_event_handler`)

- Receives `PlaybackEvent` from mixer
- Updates `SharedState`
- Emits `PlaybackPosition SSE`

|

D. Mixer Thread
(fills `AudioRingBuffer`)

- DBD-PARAM-111: Check every 10ms
- Graduated filling strategy:
 - Critical (<25%): No sleep
 - Low (25-50%): 512 frames/wake
(DBD-PARAM-112)
 - Optimal (50-75%): 256 frames/wake
(DBD-PARAM-113)
 - High (>75%): Sleep

7. CROSSFADE MIXER (`pipeline/mixer.rs`)

`CrossfadeMixer::get_next_frame()`

|
| [State Machine - SSD-MIX-010]
|
| → None (idle)
| ↳ Check `buffer_manager` for ready passages

- DBD-PARAM-088: Require 22,050 samples before start
 - ↳ Transition to Playing or Crossfading
- ↳ Playing (single passage)
 - ↳ Read frame from PlayoutRingBuffer
 - `buffer_manager.pop_frames(queue_entry_id, 1)`
 - ↳ Apply master volume
 - `frame.left *= volume`
 - `frame.right *= volume`
 - ↳ Check for crossfade trigger
 - At `fade_out_point_ticks`
 - If next passage ready → Start crossfade
 - ↳ Emit PositionUpdate events (every ~1 second)
 - `PlaybackEvent { queue_entry_id, frame_position }`
- ↳ Crossfading (two passages)
 - ↳ Read from BOTH buffers simultaneously
 - `old_frame = pop_frames(old_queue_entry_id)`
 - `new_frame = pop_frames(new_queue_entry_id)`
 - ↳ Apply crossfade curves
 - `old_weight = fade_out_curve(progress)`
 - `new_weight = fade_in_curve(progress)`
 - ↳ Mix frames
 - `mixed = (old_frame * old_weight) + (new_frame * new_weight)`
 - ↳ Apply master volume
 - ↳ When crossfade complete:
 - Transition to Playing (new passage)
 - `old_queue_entry_id` marked Finished
 - `queue.advance()`
- ↳ PausedDecay (pause mode)
 - ↳ Start from last played sample
 - ↳ Apply exponential decay per sample
 - `sample *= decay_factor`
 - DBD-PARAM-090: `pause_decay_factor = 0.96875 (31/32)`
 - ↳ Check decay floor
 - If `|sample| < pause_decay_floor` → Output 0.0
 - DBD-PARAM-100: `pause_decay_floor = 0.0001778`
 - ↳ On play: Linear fade-in over 500ms

8. AUDIO RING BUFFER (ring_buffer.rs)

AudioRingBuffer (lock-free SPSC ring buffer)

| DBD-PARAM-030: output_ringbuffer_size = 88,200 samples (2.0s)

|→ [Producer Side - Mixer Thread]

- producer.push(AudioFrame)
- Check interval: DBD-PARAM-111 = 10ms
- Graduated filling strategy:
 - Critical (<25%): Fill without sleeping
 - Low (25-50%): 512 frames (DBD-PARAM-112)
 - Optimal (50-75%): 256 frames (DBD-PARAM-113)
 - High (>75%): Sleep

|→ [Consumer Side - Audio Callback]

- consumer.pop() → Option<AudioFrame>
- Real-time thread (no locks allowed)
- Underrun detection with grace period

Mixer Thread continuously:

| DBD-PARAM-111: mixer_check_interval_ms = 10ms

|→ Check fill level

- occupied_len() / capacity()

|→ Get frames from mixer

- mixer.get_next_frame() → AudioFrame

|→ Push to ring buffer

- producer.push(frame)
- Batch size depends on fill level (256 or 512)

9. AUDIO OUTPUT (audio/output.rs:1-200)

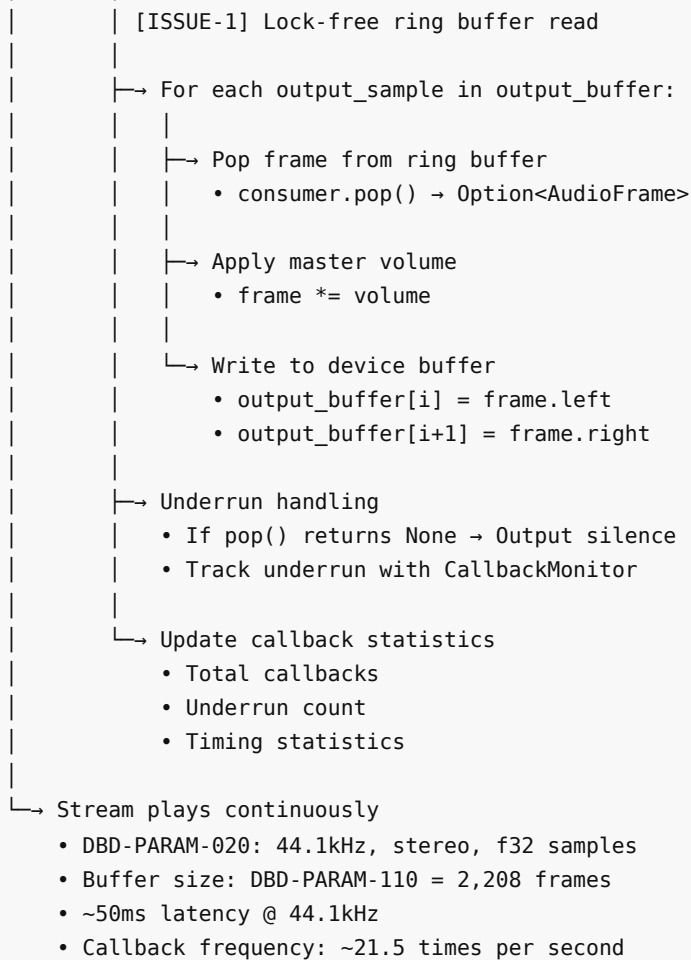
AudioOutput::start(consumer: AudioRingBufferConsumer)

| [CPAL Stream with Real-time Callback]

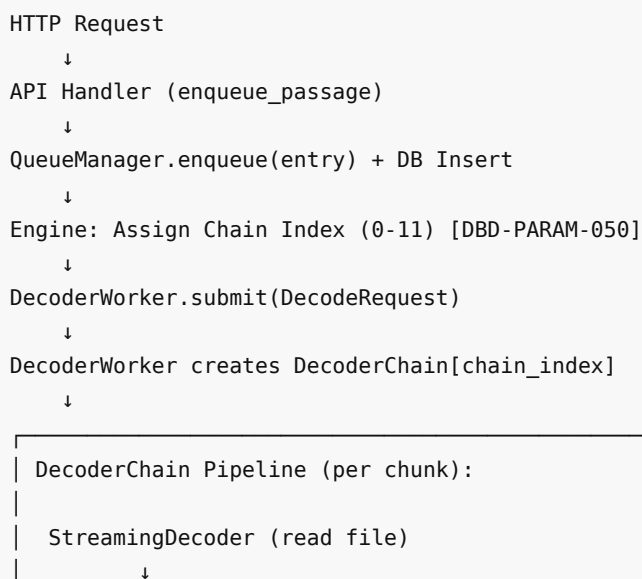
| DBD-PARAM-110: audio_buffer_size = 2,208 frames (50.1ms @ 44.1kHz)

|→ Build stream with data callback

|→ Audio Callback (runs in separate thread)



COMPLETE FLOW SUMMARY



```

| StatefulResampler (→ 44.1kHz) [DBD-PARAM-020] |
|     ↓                                         |
| Fader (apply curves)                         |
|     ↓                                         |
| PlayoutRingBuffer (via BufferManager)         |
| [DBD-PARAM-070: 661,941 samples = 15.01s]    |
| [DBD-PARAM-080: 4,410 headroom = 0.1s]       |
| [DBD-PARAM-085: 44,100 hysteresis = 1.0s]    |
|-----|
|
↓
BufferManager tracks state:
    Empty → Filling → Ready → Playing → Finished
    ↓
    Ready threshold: [DBD-PARAM-088] 22,050 samples (0.5s)
    ↓
    Emits ReadyForStart event
    ↓
Buffer Event Handler notifies Mixer
    ↓
CrossfadeMixer reads frames from PlayoutRingBuffer:
    • Single passage: Playing mode
    • Two passages: Crossfading mode
    • Pause: Exponential decay [DBD-PARAM-090, DBD-PARAM-100]
    ↓
Mixer Thread pushes AudioFrame to AudioRingBuffer
    [DBD-PARAM-030: 88,200 samples = 2.0s]
    [DBD-PARAM-111: Check every 10ms]
    [DBD-PARAM-112: 512 frames when low]
    [DBD-PARAM-113: 256 frames when optimal]
    ↓
Audio Callback (real-time thread) pops from AudioRingBuffer
    [DBD-PARAM-110: 2,208 frames/callback = 50.1ms]
    ↓
CPAL Stream → Audio Device → Speakers
    [DBD-PARAM-020: 44,100 Hz stereo]

```

KEY CONCURRENCY NOTES

1. Single-threaded Decoder Worker (DBD-DEC-040)
 - Processes chains serially for cache coherency
 - Maintains HashMap of active DecoderChains
 - DBD-PARAM-050: Up to 12 chains (configurable)
 - DBD-PARAM-060: Priority check every 5000ms
2. Lock-free Ring Buffers (ISSUE-1)
 - PlayoutRingBuffer: Decoder → Mixer (per passage)
 - DBD-PARAM-070: 661,941 samples capacity
 - AudioRingBuffer: Mixer → Audio Callback (output)
 - DBD-PARAM-030: 88,200 samples capacity

- No mutexes in audio callback path

3. Event-driven Architecture

- BufferEvent channel: BufferManager → Engine
- PlaybackEvent channel: Mixer → Position Handler
- SSE broadcast: SharedState → HTTP clients

4. Graduated Filling Strategy (mixer thread)

- DBD-PARAM-111: Check interval 10ms
- DBD-PARAM-112: 512 frames when buffer <50%
- DBD-PARAM-113: 256 frames when buffer 50-75%
- Prevents underruns with adaptive batch sizes

5. Chain Assignment Pool (DBD-LIFECYCLE-030/040)

- DBD-PARAM-050: 0-11 chain indices (default 12 chains)
- Lowest-numbered allocation strategy
- Persistent queue_entry_id → chain_index mapping

TRACEABILITY REFERENCES

- SSD-FLOW-010: Complete playback sequence
- DBD-BUF-010 through DBD-BUF-080: Buffer lifecycle management
- DBD-DEC-040, DBD-DEC-090, DBD-DEC-110: Decoder architecture
- DBD-PARAM-020 through DBD-PARAM-113: Configurable parameters (see next section)
- PERF-POLL-010: Event-driven buffer readiness (instant startup)
- ISSUE-1: Lock-free audio callback using ring buffer
- REQ-AP-ERR-010/011: Error handling and event emission

5. DBD-PARAM Parameter Mapping

Parameter Overview Table

Parameter	Name	Default	Unit	Applied Where
DBD-PARAM-010	General	N/A	-	Settings table storage
DBD-PARAM-020	working_sample_rate	44,100	Hz	Resampler output, throughout pipeline
DBD-PARAM-030	output_ringbuffer_size	88,200	samples	AudioRingBuffer capacity (2.0s)
DBD-PARAM-040	output_refill_period	90	ms	Mixer check interval (deprecated)

DBD-PARAM-050	maximum_decode_streams	12	count	Chain pool size, max concurrent decoders
DBD-PARAM-060	decode_work_period	5,000	ms	Priority queue check interval
DBD-PARAM-065	decode_chunk_size	25,000	samples	Resampler output per chunk (~1s @ 44.1kHz)
DBD-PARAM-070	playout_ringbuffer_size	661,941	samples	PlayoutRingBuffer capacity (15.01s)
DBD-PARAM-080	playout_ringbuffer_headroom	4,410	samples	Decoder pause threshold (0.1s)
DBD-PARAM-085	decoder_resume_hysteresis	44,100	samples	Pause/resume gap (1.0s)
DBD-PARAM-088	mixer_min_start_level	22,050	samples	Buffer ready threshold (0.5s)
DBD-PARAM-090	pause_decay_factor	0.96875	ratio	Exponential decay per sample in pause mode
DBD-PARAM-100	pause_decay_floor	0.0001778	level	Minimum level before outputting zero
DBD-PARAM-110	audio_buffer_size	2,208	frames	CPAL callback buffer size (50.1ms)
DBD-PARAM-111	mixer_check_interval_ms	10	ms	Mixer thread wake frequency
DBD-PARAM-112	mixer_batch_size_low	512	frames	Frames filled when buffer <50%
DBD-PARAM-113	mixer_batch_size_optimal	256	frames	Frames filled when buffer 50-75%
DBD-PARAM-	Batch size rationale	-	-	Design documentation

114				
DBD- PARAM- 120	Default value rationale	-	-	Design documentation

Visual Parameter Mapping

```

graph TD
    subgraph "Chain Pool [DBD-PARAM-050]"
        CP[Maximum Decode Streams<br/>default: 12 chains<br/>indices 0-11]
    end

    subgraph "DecoderChain Pipeline"
        D[StreamingDecoder]
        R[StatefulResampler<br/>[DBD-PARAM-020]<br/>working_sample_rate<br/>44,100 Hz]
        F[Fader]

        D -->|chunks| R
        R -->|[DBD-PARAM-065]<br/>decode_chunk_size<br/>25,000 samples/chunk| F
    end

    subgraph "PlayoutRingBuffer [DBD-PARAM-070/080/085]"
        PRB[Capacity: 661,941 samples<br/>[DBD-PARAM-070]<br/>15.01 seconds @ 44.1kHz]

        Pause[Decoder Pause Logic]
        Resume[Decoder Resume Logic]

        Pause -->|free_space ≤ 4,410<br/>[DBD-PARAM-080]| Pause
        Resume -->|free_space ≥ 48,510<br/>[DBD-PARAM-085]+[080]| Resume
    end

    subgraph "BufferManager State Machine"
        Empty[Empty]
        Filling[Filling]
        Ready[Ready<br/>[DBD-PARAM-088]<br/>mixer_min_start_level<br/>22,050 samples]
        Playing[Playing]

        Empty --> Filling
        Filling --> Ready
        Ready --> Playing
    end

    subgraph "CrossfadeMixer Pause Mode [DBD-PARAM-090/100]"
        PD[PausedDecay<br/>sample *= 0.96875<br/>[DBD-PARAM-090]]
        PS[PausedSilence<br/>level < 0.0001778<br/>[DBD-PARAM-100]]

        PD -->|decay per sample| PS
    end

```

```

end

subgraph "Mixer Thread [DBD-PARAM-111/112/113]"
    MT[Check Interval<br/>10 ms<br/>[DBD-PARAM-111]]

    Low[Buffer <50%<br/>Fill 512 frames<br/>[DBD-PARAM-112]]
    Opt[Buffer 50-75%<br/>Fill 256 frames<br/>[DBD-PARAM-113]]

    MT --> Low
    MT --> Opt
end

subgraph "AudioRingBuffer [DBD-PARAM-030]"
    ARB[Capacity: 88,200 samples<br/>2.0 seconds @ 44.1kHz]
end

subgraph "Audio Callback [DBD-PARAM-110]"
    AC[Buffer Size: 2,208 frames<br/>50.1 ms @ 44.1kHz<br/>~21.5
callbacks/second]
end

subgraph "Priority Queue [DBD-PARAM-060]"
    PQ[Check every 5,000 ms<br/>Immediate > Next > Prefetch]
end

CP --> D
F --> PRB
PRB --> Ready
Ready --> PD
Playing --> Low
Playing --> Opt
Low --> ARB
Opt --> ARB
ARB --> AC
PQ --> D

style CP fill:#fff3e0
style R fill:#e8f5e9
style PRB fill:#e1f5ff
style Ready fill:#fff9c4
style PD fill:#fce4ec
style MT fill:#f3e5f5
style ARB fill:#e0f2f1
style AC fill:#e8eaf6

```

Parameter Application Points

1. Chain Pool Configuration

- **DBD-PARAM-050** (`maximum_decode_streams` = 12)
 - **Applied in:** `engine.rs:289` - Initialize available chains pool
 - **Purpose:** Limits concurrent decoder chains

- **Memory impact:** Each chain allocates a `PlayoutRingBuffer` (~5.3 MB)
- **Total memory:** 12 chains × 5.3 MB = ~64 MB for playout buffers

2. Decoder Pipeline

- **DBD-PARAM-020** (`working_sample_rate` = 44,100 Hz)
 - **Applied in:** `decoder_chain.rs:148` - `StatefulResampler` target rate
 - **Purpose:** Standardizes all audio to 44.1kHz for mixing
 - **Affects:** All downstream components (fader, mixer, output)
- **DBD-PARAM-065** (`decode_chunk_size` = 25,000 samples)
 - **Applied in:** `decoder_chain.rs:145` - Chunk size calculation
 - **Purpose:** Balances decode granularity vs overhead
 - **Equivalent:** ~566ms of audio @ 44.1kHz per chunk
- **DBD-PARAM-060** (`decode_work_period` = 5,000 ms)
 - **Applied in:** `decoder_worker.rs:190` - Priority queue check interval
 - **Purpose:** Prevents low-priority long decodes from starving high-priority
 - **Behavior:** Check priority between chunks, yield if higher priority pending

3. PlayoutRingBuffer Configuration

- **DBD-PARAM-070** (`playout_ringbuffer_size` = 661,941 samples)
 - **Applied in:** `playout_ring_buffer.rs:126` - Buffer capacity
 - **Purpose:** Holds decoded audio per passage
 - **Equivalent:** 15.01 seconds @ 44.1kHz stereo
 - **Memory:** $661,941 \times 8 \text{ bytes (f32 stereo)} = \sim 5.3 \text{ MB per buffer}$
- **DBD-PARAM-080** (`playout_ringbuffer_headroom` = 4,410 samples)
 - **Applied in:** `buffer_manager.rs:264` - Decoder pause threshold
 - **Purpose:** Reserves space for in-flight resampler output
 - **Equivalent:** 0.1 seconds @ 44.1kHz
 - **Behavior:** Decoder pauses when `free_space` ≤ 4,410
- **DBD-PARAM-085** (`decoder_resume_hysteresis` = 44,100 samples)
 - **Applied in:** `buffer_manager.rs:124` - Hysteresis configuration
 - **Purpose:** Prevents pause/resume oscillation
 - **Equivalent:** 1.0 second gap @ 44.1kHz
 - **Resume threshold:** $44,100 + 4,410 = 48,510$ samples free space

4. Buffer State Machine

- **DBD-PARAM-088** (`mixer_min_start_level` = 22,050 samples)
 - **Applied in:** `buffer_manager.rs:282` - Ready threshold calculation
 - **Purpose:** Ensures sufficient buffer before playback starts
 - **Equivalent:** 0.5 seconds @ 44.1kHz
 - **First passage optimization:** May use 500ms instead of 3000ms
 - **Event:** Emits `BufferEvent::ReadyForStart` when threshold reached

5. Mixer Pause Mode

- **DBD-PARAM-090** (`pause_decay_factor` = 0.96875)

- **Applied in:** `mixer.rs` (pause decay logic)
- **Purpose:** Exponential decay to zero in pause mode
- **Formula:** `sample *= 0.96875` recursively each sample
- **Effect:** Reduces "pop" when pausing
- **DBD-PARAM-100** (`pause_decay_floor` = 0.0001778)
 - **Applied in:** `mixer.rs` (pause floor check)
 - **Purpose:** Threshold for switching to silence
 - **Behavior:** When `|sample| < 0.0001778`, output 0.0 instead

6. Mixer Thread Configuration

- **DBD-PARAM-111** (`mixer_check_interval_ms` = 10 ms)
 - **Applied in:** `engine.rs:421` - Mixer thread wake interval
 - **Purpose:** Controls mixer thread frequency
 - **Trade-off:** 10ms = lower CPU, more stable vs 5ms = more responsive
- **DBD-PARAM-112** (`mixer_batch_size_low` = 512 frames)
 - **Applied in:** `engine.rs:488` - Low fill level batch size
 - **Purpose:** Aggressive recovery when buffer <50%
 - **Behavior:** Fill 512 frames per wake when critically low
- **DBD-PARAM-113** (`mixer_batch_size_optimal` = 256 frames)
 - **Applied in:** `engine.rs:500` - Optimal fill level batch size
 - **Purpose:** Steady-state operation when buffer 50-75%
 - **Behavior:** Fill 256 frames per wake to maintain level

7. Output Ring Buffer

- **DBD-PARAM-030** (`output_ringbuffer_size` = 88,200 samples)
 - **Applied in:** `ring_buffer.rs` - AudioRingBuffer capacity
 - **Purpose:** Lock-free buffer between mixer and audio callback
 - **Equivalent:** 2.0 seconds @ 44.1kHz stereo
 - **Memory:** $88,200 \times 8 \text{ bytes} = \sim 706 \text{ KB}$

8. Audio Output

- **DBD-PARAM-110** (`audio_buffer_size` = 2,208 frames)
 - **Applied in:** `output.rs:144` - CPAL StreamConfig buffer size
 - **Purpose:** Audio device callback buffer size
 - **Equivalent:** 50.1 ms @ 44.1kHz
 - **Callback frequency:** $44,100 / 2,208 = \sim 21.5$ times per second
 - **Latency impact:** Direct contribution to output latency

Parameter Interdependencies

```
graph LR
  subgraph "Time Domain Parameters"
    T20[DBD-PARAM-020<br/>working_sample_rate<br/>44,100 Hz]
    T65[DBD-PARAM-065<br/>decode_chunk_size<br/>25,000 samples]
    T70[DBD-PARAM-070<br/>playout_ringbuffer_size<br/>661,941 samples]
    T88[DBD-PARAM-088<br/>mixer_min_start_level<br/>22,050 samples]
```

```

    T30[DBD-PARAM-030<br/>output_ringbuffer_size<br/>88,200 samples]
    T110[DBD-PARAM-110<br/>audio_buffer_size<br/>2,208 frames]
end

subgraph "Interval Parameters"
    I60[DBD-PARAM-060<br/>decode_work_period<br/>5,000 ms]
    I111[DBD-PARAM-111<br/>mixer_check_interval<br/>10 ms]
end

subgraph "Threshold Parameters"
    TH80[DBD-PARAM-080<br/>headroom<br/>4,410 samples]
    TH85[DBD-PARAM-085<br/>hysteresis<br/>44,100 samples]
    TH90[DBD-PARAM-090<br/>decay_factor<br/>0.96875]
    TH100[DBD-PARAM-100<br/>decay_floor<br/>0.0001778]
end

subgraph "Batch Parameters"
    B112[DBD-PARAM-112<br/>batch_size_low<br/>512 frames]
    B113[DBD-PARAM-113<br/>batch_size_optimal<br/>256 frames]
end

T20 --> T65
T20 --> T70
T20 --> T88
T20 --> T30
T20 --> T110

I111 --> B112
I111 --> B113

T110 -.fill rate.-> B112
T110 -.fill rate.-> B113

TH80 --> TH85

style T20 fill:#e8f5e9
style I111 fill:#f3e5f5
style TH80 fill:#fff9c4
style B112 fill:#e1f5ff

```

Configuration Access Pattern

All DBD-PARAM parameters are:

1. **Stored:** SQLite `settings` table (see IMPL001 Database Schema)
2. **Loaded:** Once at startup in `engine.rs:173-194`
3. **Applied:** Throughout component initialization
4. **Restart required:** Changes require application restart

Example: Loading Parameters at Startup

```
// engine.rs:173-194 - Parallel configuration loading
let (initial_volume, min_buffer_threshold, interval_ms, grace_period_ms,
    mixer_config, maximum_decode_streams, resume_hysteresis,
    mixer_min_start_level, audio_buffer_size) = tokio::join!(
    db::settings::get_volume(&db_pool), // Volume level
    db::settings::load_minimum_buffer_threshold(&db_pool), // DBD-PARAM-088
    db::settings::load_position_event_interval(&db_pool), // Position SSE
    db::settings::load_ring_buffer_grace_period(&db_pool), // Underrun grace
    db::settings::load_mixer_thread_config(&db_pool), // DBD-PARAM-
111/112/113
    db::settings::load_maximum_decode_streams(&db_pool), // DBD-PARAM-050
    db::settings::get_decoder_resume_hysteresis(&db_pool), // DBD-PARAM-085
    db::settings::load_mixer_min_start_level(&db_pool), // DBD-PARAM-088
    db::settings::load_audio_buffer_size(&db_pool), // DBD-PARAM-110
);
```

This parallel loading pattern (PERF-INIT-010) reduces startup time from ~45ms (sequential) to ~10ms (parallel).

References

Primary Documentation

- [SPEC016 Decoder Buffer Design](#) - Complete parameter definitions
- [SPEC013 Single Stream Playback](#) - Architecture overview
- [SPEC014 Single Stream Design](#) - Design details
- [SPEC002 Crossfade Design](#) - Crossfading mechanics

Implementation Files

- `engine.rs` - Playback engine coordination
- `decoder_worker.rs` - Single-threaded decoder loop
- `decoder_chain.rs` - Decode → resample → fade → buffer pipeline
- `buffer_manager.rs` - Buffer lifecycle state machine
- `pipeline/mixer.rs` - Crossfade mixer
- `audio/output.rs` - CPAL audio output
- `ring_buffer.rs` - Lock-free AudioRingBuffer
- `playout_ring_buffer.rs` - Lock-free PlayoutRingBuffer

Traceability

- **SSD-FLOW-010** - Complete playback sequence
- **DBD-BUF-010 through DBD-BUF-080** - Buffer lifecycle management
- **DBD-DEC-040, DBD-DEC-090, DBD-DEC-110** - Decoder architecture
- **DBD-LIFECYCLE-010 through DBD-LIFECYCLE-060** - Chain assignment lifecycle
- **PERF-POLL-010** - Event-driven buffer readiness
- **ISSUE-1** - Lock-free audio callback

DOCUMENT STATUS: Complete

LAST UPDATED: 2025-10-27

MAINTAINED BY: Development team